Evolutionary Computation (Lab 4: Local Search with Candidates)

Oliwier Necelman 160300 Karol Okoń 160323

Problem story

The student still has to attend 100 of the 200 classes in one day while getting as little tired as possible. This week they do not build the plan from scratch, but take a ready random plan and keep improving it with small changes: replacing one attended class with one not attended, or walking a part of the route in reverse. To plan faster, they now consider only changes that add a walk from a class to one of its 10 most convenient classes (short walk, not tiring class).

Problem description (exactly the same as last week)

We are given a set of n nodes. Each node has x and y coordinates in the plane and a cost. The goal is to select exactly $k = \lceil n/2 \rceil$ nodes and form a Hamiltonian cycle through them that minimizes the total length of the cycle plus the total cost of the selected nodes:

$$f(C) = \sum_{i=1}^{k} d\left(C_i, C_{(i \bmod k)+1}\right) + \sum_{i=1}^{k} c(C_i),$$

where $C = (C_1, \ldots, C_k)$ is the cycle, d is the distance between two nodes and c is the cost of a node.

Distances are Euclidean distances rounded mathematically to integers. The distance matrix is computed once, right after reading an instance. The optimization methods access only the distance matrix and the node costs, never the coordinates, so the same code works for instances defined only by a distance matrix, including non-Euclidean distances for which the triangle inequality does not hold.

Both instances used in the experiments (A and B) have $n = 200$ nodes, so every solution contains $k = 100$ nodes. Node costs range from 0 to 1997 in instance A and from 3 to 999 in instance B.

Implemented algorithms

Notation: D is the distance matrix, c the vector of node costs, C the current cycle with k nodes, C_i its node at position i (indices modulo k), and U the set of nodes not in the cycle.

Neighborhood

The neighborhood is the best one from lab 3. Every move is evaluated by its delta, the change of the objective function, computed only from the edges and nodes that change; a negative delta means an improvement. The neighbor solution is never built explicitly and the objective is never recomputed from scratch.

- **Inter-route move**, exchange of a selected and an unselected node: node C_i is replaced by $u \in U$. With $a = C_{i-1}$, $v = C_i$, $b = C_{i+1}$:

$$\Delta = D[a][u] + D[u][b] + c[u] - D[a][v] - D[v][b] - c[v].$$

- **Intra-route move, two-edges exchange**: edges (C_i, C_{i+1}) and (C_j, C_{j+1}) are replaced by (C_i, C_j) and (C_{i+1}, C_{j+1}), which reverses the part of the cycle between them:

$$\Delta = D[C_i][C_j] + D[C_{i+1}][C_{j+1}] - D[C_i][C_{i+1}] - D[C_j][C_{j+1}].$$

  Pairs of edges sharing a node are skipped, as the move changes nothing.

Candidate moves

For each node v we determine the $m = 10$ "nearest" other nodes, where "nearest" means the smallest value of $D[v][u] + c[u]$ (edge length plus the cost of the other node). An edge is a **candidate edge** if at least one of its endpoints chose the other one, so the lists are made symmetric; after this a list contains on average 17.7 nodes on instance A and 16.2 nodes on instance B. A move is a **candidate move** if it adds at least one candidate edge to the solution. For a candidate edge (v, u) with $v = C_i$ in the cycle there are two such moves:

- u **not in the cycle (inter-route move):** u replaces the predecessor C_{i-1} or the successor C_{i+1} of v. Replacing v itself would not add the edge (v, u).
- $u = C_j$ **in the cycle (two-edges exchange):** either the edges (C_i, C_{i+1}) and (C_j, C_{j+1}), or the edges (C_{i-1}, C_i) and (C_{j-1}, C_j) are removed; both moves add the edge (C_i, C_j).

The candidate lists depend only on the instance. In our implementation they are built at the beginning of every run, and this time is included in the measured running time.

Local search with candidate moves

Steepest local search starting from a random solution. The baseline is the same steepest local search (two-edges exchange, random start) from lab 3, which evaluates the whole neighborhood in every step.

```
Input:  distance matrix D, node costs c, seed s, number of candidates m = 10
Output: locally optimal cycle C

C ← random solution (seed s)
for each node v:
    N[v] ← m nodes u ≠ v with the smallest D[v][u] + c[u]
make the lists symmetric: if u ∈ N[v], add v to N[u]
repeat
    improved ← steepest step(C)
until not improved
return C
```

Steepest step with candidate moves — the loops go only over the candidate edges, so only candidate moves are evaluated:

```
best ← none, Δbest ← 0
for each node v in C:
    for each u ∈ N[v]:                              ▷ candidate edge (v, u)
        if u ∉ C:                                   ▷ inter-route moves
            for p in (position before v, position after v):
                if Δinter(p, u) < Δbest:
                    best ← (inter, p, u), Δbest ← Δinter(p, u)
        else if v < u:                              ▷ two-edges exchange
            i ← position of v, j ← position of u
            for (p, q) in ((i, j), (i − 1, j − 1)):
                if Δ2edges(p, q) < Δbest:
                    best ← (2-edges, p, q), Δbest ← Δ2edges(p, q)
if best = none: return false                        ▷ local optimum
apply best, return true
```

The condition $v < u$ prevents evaluating the same pair of nodes twice. A move can still be evaluated twice when it adds two candidate edges at once; on random starting solutions this concerns about 5–6% of the evaluated

moves. Instead of all 14 950 moves of the full neighborhood (10 000 inter-route moves and 4 950 pairs of edges), a step on a random starting solution evaluates only about 2 400–2 700 moves.

Starting solutions

- **Random start**: a random set of k nodes in random order. Both methods start from the same 200 random solutions.

Computational experiment

Both methods of this assignment, steepest local search with two-edges exchange from a random start without candidate moves (the baseline) and with candidate moves, were run 200 times per instance from the same 200 random starting solutions (seeds 1–200). For reference, the tables also contain all earlier methods, run again in the same experiment. The construction heuristics give the same objective values as in the lab 1 and lab 2 reports. The random solutions differ from the earlier reports, because the random number generator is now seeded inside the algorithm, so the random solution and local search from a random start differ slightly from the lab 3 report. Local search with a greedy start now uses greedy cycle with weighted 2-regret on both instances (in lab 3, nearest neighbor II with weighted 2-regret was used on instance B).

As in the previous labs, each run is a separate execution of the `test_algo` program (g++ with `-O2`), and only the time of the `solve` call is measured. For local search this time includes building the random starting solution and, with candidate moves, building the candidate lists. All values in the tables come from one run of the experiment.

Objective function

Method	Instance A	Instance B
	av (min – max)	av (min – max)
Random solution	262,653 (238,917 – 286,724)	206,842 (181,610 – 226,734)
NearestNeighborI	87,477 (84,362 – 90,075)	59,397 (56,007 – 63,166)
NearestNeighborII	74,307 (72,328 – 77,022)	59,476 (52,915 – 66,164)
GreedyCycle	73,739 (72,590 – 75,281)	59,651 (55,655 – 63,734)
Nearest neighbor II 2-regret	122,861 (111,527 – 133,181)	74,763 (69,032 – 80,767)
Nearest neighbor II weighted 2-regret	75,561 (72,419 – 78,504)	55,698 (50,485 – 61,406)
Greedy cycle 2-regret	120,995 (112,654 – 130,391)	73,459 (69,609 – 78,175)
Greedy cycle weighted 2-regret	73,353 (72,011 – 75,532)	57,863 (52,731 – 63,370)
LS greedy, two-nodes, random start	88,669 (82,989 – 96,524)	64,576 (57,464 – 71,404)
LS greedy, two-nodes, greedy start	73,022 (71,955 – 74,861)	57,528 (52,702 – 62,933)
LS greedy, two-edges, random start	75,557 (73,618 – 77,679)	52,452 (50,399 – 55,468)
LS greedy, two-edges, greedy start	72,966 (71,955 – 74,861)	57,196 (52,542 – 62,428)
LS steepest, two-nodes, random start	89,321 (79,649 – 96,354)	65,871 (59,357 – 74,662)
LS steepest, two-nodes, greedy start	73,022 (71,955 – 74,861)	57,525 (52,702 – 62,933)
LS steepest, two-edges, random start	76,161 (74,008 – 80,204)	52,570 (50,492 – 56,091)
LS steepest, two-edges, greedy start	73,004 (71,955 – 74,861)	57,126 (52,542 – 62,301)
LS with candidates	76,234 (73,878 – 79,182)	52,546 (50,422 – 55,822)

Running time [ms]

Method	Instance A	Instance B
	av (min – max)	av (min – max)
Random solution	0.004 (0.003 – 0.009)	0.003 (0.003 – 0.009)
NearestNeighborI	0.029 (0.026 – 0.049)	0.029 (0.027 – 0.050)
NearestNeighborII	1.130 (1.083 – 1.923)	1.157 (1.082 – 2.056)
GreedyCycle	0.778 (0.755 – 1.221)	0.792 (0.755 – 1.357)
Nearest neighbor II 2-regret	1.008 (0.940 – 1.572)	1.054 (0.933 – 1.657)
Nearest neighbor II weighted 2-regret	1.037 (0.922 – 1.661)	1.036 (0.883 – 1.519)
Greedy cycle 2-regret	1.365 (1.286 – 1.902)	1.379 (1.288 – 1.831)
Greedy cycle weighted 2-regret	1.338 (1.262 – 1.950)	1.311 (1.219 – 1.815)
LS greedy, two-nodes, random start	3.589 (1.732 – 5.992)	2.437 (1.115 – 4.002)
LS greedy, two-nodes, greedy start	1.600 (1.382 – 3.926)	1.642 (1.404 – 2.435)
LS greedy, two-edges, random start	2.964 (1.431 – 4.742)	2.174 (1.046 – 4.796)
LS greedy, two-edges, greedy start	1.559 (1.361 – 2.122)	1.732 (1.430 – 2.684)
LS steepest, two-nodes, random start	10.672 (8.225 – 17.996)	11.100 (8.556 – 20.230)
LS steepest, two-nodes, greedy start	1.609 (1.349 – 2.072)	1.758 (1.393 – 3.084)
LS steepest, two-edges, random start	7.029 (6.200 – 11.413)	7.200 (6.120 – 10.756)
LS steepest, two-edges, greedy start	1.557 (1.331 – 2.676)	1.879 (1.442 – 3.143)
LS with candidates	3.540 (3.213 – 4.673)	3.325 (2.870 – 4.806)

Local search with candidate moves reaches practically the same objective values as the baseline (LS steepest, two-edges, random start) in about half the time: 3.54 ms instead of 7.03 ms on instance A (2.0 times faster) and 3.33 ms instead of 7.20 ms on instance B (2.2 times faster).

The next table shows the number of steps, i.e. improving moves applied before a local optimum was reached. Zero steps mean that the starting solution (a construction from a greedy start) was already a local optimum.

Number of steps

Method	Instance A	Instance B
	av (min – max)	av (min – max)
LS greedy, two-nodes, random start	525.0 (412 – 605)	521.4 (445 – 603)
LS greedy, two-nodes, greedy start	4.3 (0 – 17)	5.8 (1 – 32)
LS greedy, two-edges, random start	510.2 (449 – 578)	510.5 (459 – 596)
LS greedy, two-edges, greedy start	5.2 (0 – 19)	10.2 (2 – 39)
LS steepest, two-nodes, random start	154.1 (123 – 201)	151.4 (123 – 200)
LS steepest, two-nodes, greedy start	3.3 (0 – 9)	5.3 (1 – 21)
LS steepest, two-edges, random start	133.2 (119 – 150)	134.0 (119 – 152)
LS steepest, two-edges, greedy start	3.3 (0 – 10)	9.2 (2 – 24)
LS with candidates	132.7 (116 – 154)	132.5 (115 – 150)

"2-regret" denotes $\alpha = 1$ and "weighted 2-regret" denotes $\alpha = 0.5$. "LS greedy" and "LS steepest" are the two types of local search, "two-nodes" and "two-edges" the intra-route move, and "random" and "greedy" the starting solution.

Visualization of the best solutions

Each panel shows the best solution of the baseline (steepest local search with two-edges exchange from a random start, left) and of local search with candidate moves (right). Node color encodes the node cost (darker means more expensive). Selected nodes are drawn larger, the solid line shows the cycle, and the dotted line marks the edge from the last node back to the first.

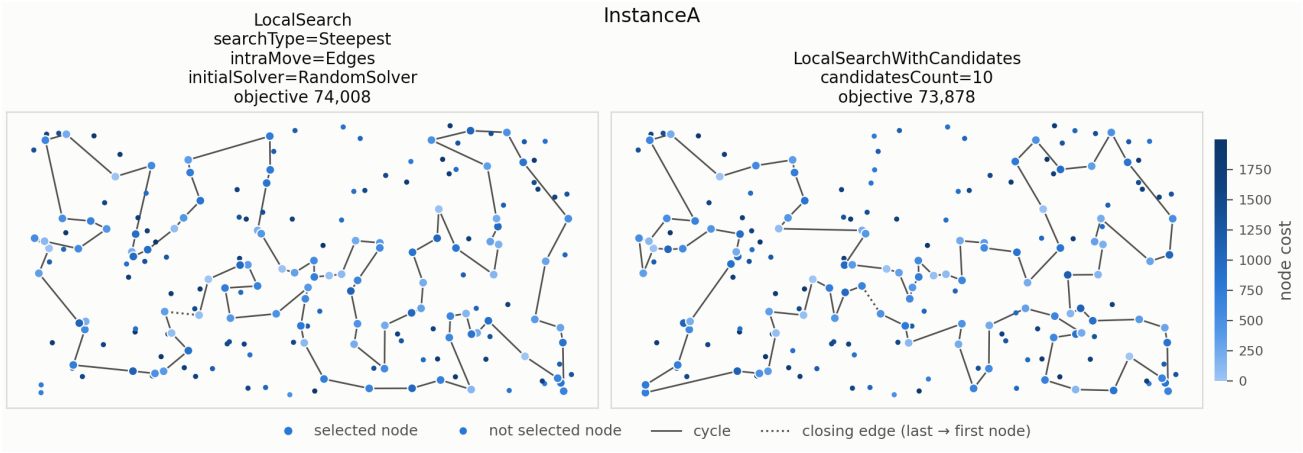


Figure 1: Best solutions for instance A

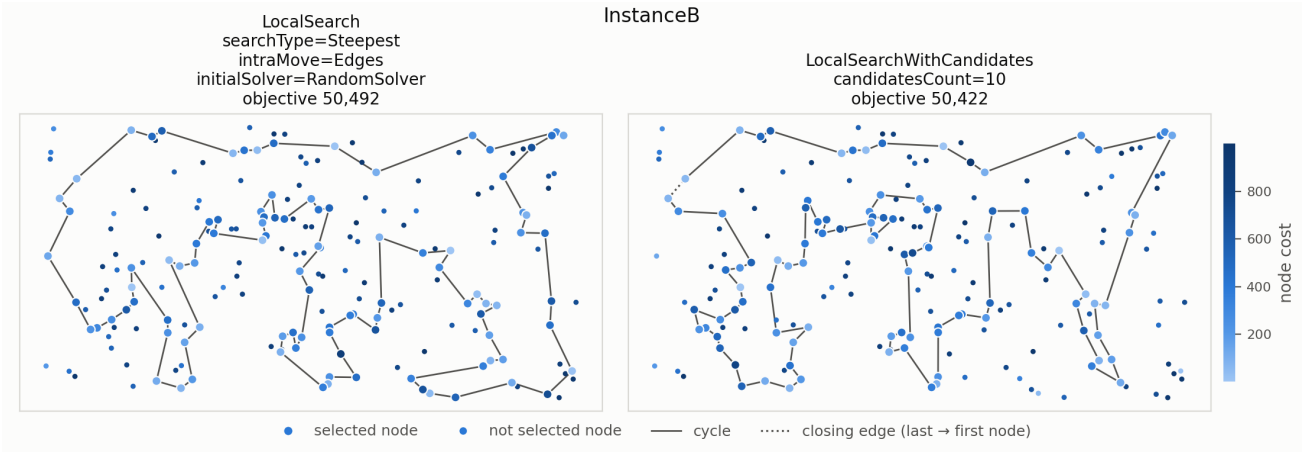


Figure 2: Best solutions for instance B

Best solutions

The best solutions of the earlier methods are listed in the lab 1 and lab 2 reports.

Instance A

LS steepest, two-edges, random start (objective 74 008)

30 187 189 2 171 92 69 127 97 95 168 63 116 149 49 152 108 60 31 48 80 91 100 123 83 90 132 76 4 71 157 115 112 139 51 16 67 144

Local search with candidates (objective 73 878)

185 102 87 28 180 141 30 187 2 171 92 104 129 69 127 95 97 160 83 123 100 152 49 11 63 168 116 149 60 108 31 48 80 91 157 71 4 76

Instance B

LS steepest, two-edges, random start (objective 50 492)

30 38 47 140 154 40 60 45 134 175 71 155 125 170 115 180 6 29 132 178 65 145 28 149 116 147 42 121 198 96 48 89 57 190 70 94 105

Local search with candidates (objective 50 422)

149 116 147 42 121 198 96 48 89 83 57 190 70 81 94 105 196 114 10 90 129 153 14 2 74 20 182 80 127 54 192 162 104 185 76 66 67 24

Effect of candidates count on local search

The candidate count m controls the size of the neighborhood considered by the local search. These runs were made separately from the main experiment, so the times differ slightly from the tables above (e.g. about 4.2 ms instead of 3.54 ms for 10 candidates on instance A). Increasing it from 3 to 10 substantially improves the average objective value on both instances: from 78,779 to 76,234 on A (about 3.2%) and from 53,728 to 52,546 on B (about 2.2%). Further increases bring only small gains. On A, the best average in this experiment is 76,152 at a count of 50, only about 0.1% below the result at 10. On B, the best average is 52,529 at a count of 8; counts from 8 to 50 remain close, between 52,529 and 52,628.

Runtime, in contrast, rises steadily with the candidate count. From 3 to 50 it increases from 2.22 ms to 14.32 ms on A and from 2.00 ms to 12.17 ms on B. The average number of improving steps changes much less: it falls from about 143 to 133 on A and stays around 133–134 on B. Thus, the extra runtime mainly comes from evaluating a larger neighborhood, not from taking substantially more search steps.

Overall, a count around 8–10 offers a good quality/runtime trade-off for these instances: most of the objective improvement is already achieved, while runtime remains far below that at counts of 30 or 50. The best setting may depend on whether a small additional quality gain is worth the extra runtime.

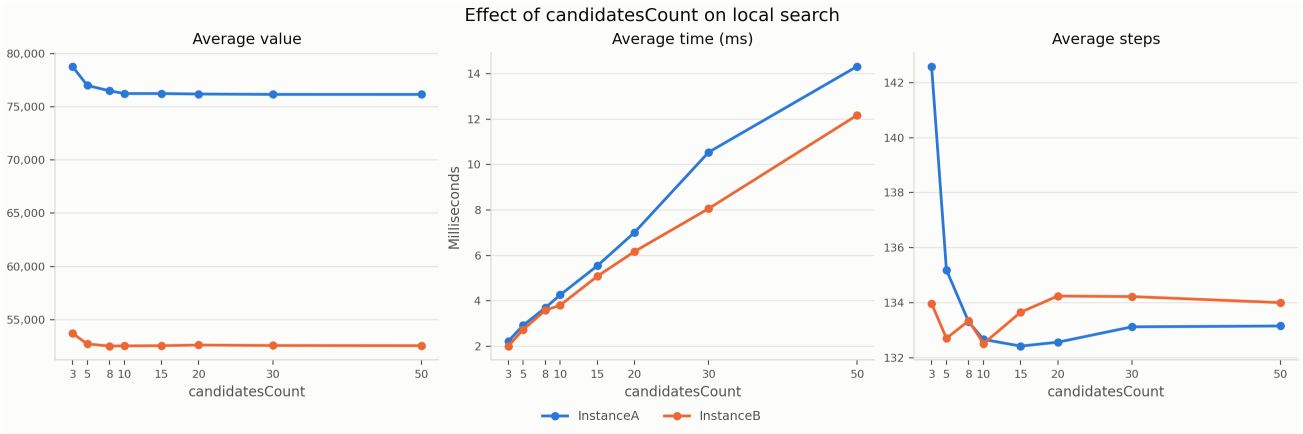


Figure 3: Effect of candidatesCount on average score, time, and steps

Solution checker

The best solutions have been checked with the Solution checker.xlsx provided on ekursy, and we confirmed that the values computed by the checker match those computed by our algorithms.

Conclusions

- **Candidate moves halve the running time without loss of quality.** Local search with candidate moves is 2.0 times faster on instance A (3.54 ms vs 7.03 ms) and 2.2 times faster on instance B (3.33 ms vs 7.20 ms). Its average objective is practically the same: 76 234 vs 76 161 on instance A (0.1% worse) and 52 546 vs 52 570 on instance B (slightly better).
- **The best and the worst runs are even slightly better with candidate moves**: best 73 878 vs 74 008 and 50 422 vs 50 492, worst 79 182 vs 80 204 and 55 822 vs 56 091. Restricting the search to candidate moves skips almost no good moves: improving moves nearly always add short edges to cheap nodes.
- **The number of steps does not change** (132.7 vs 133.2 on instance A, 132.5 vs 134.0 on instance B). Candidate moves reduce the work in each step, not the number of improvements, so one step takes about 0.026 ms instead of 0.053 ms.
- **The speedup is smaller than the reduction of evaluated moves.** A step evaluates about 6 times fewer moves, but the run is only about 2 times faster. Part of the difference is building the candidate lists

in every run (about 1.3 ms when measured separately on one of our computers), which could be done only once per instance.

- **The running time is more stable with candidate moves**: 3.21–4.67 ms vs 6.20–11.41 ms on instance A, and 2.87–4.81 ms vs 6.12–10.76 ms on instance B.
- **About 10 candidates is a good choice.** With 3 candidates the average is 3.2% (A) and 2.2% (B) worse, while more than 10 candidates improve it by at most 0.1% and cost a lot of time: with 20 candidates the time on instance A is about the same as without candidate moves, and with 30 and 50 candidates the method is slower than the baseline on both instances.
- **Greedy local search without candidate moves is still faster.** From a random start, greedy local search with two-edges exchange takes 2.96 ms (A) and 2.17 ms (B) and gives similar averages (75 557 and 52 452), so applying candidate moves to the greedy version could speed it up further.

Source code

The source code is available at https://github.com/Kari3600/EC-Project.

The local search with candidate moves is in `src/lab4_algos/`, and a single run can be executed with `src/TestAlgo.cpp`, e.g. `test_algo LocalSearchWithCandidates data/InstanceA.csv --seed 1 --candidatesCount 10`; the baseline is `LocalSearch` with `--searchType Steepest --intraMove Edges --initialSolver RandomSolver`. The experiment is run by `scripts/run_experiment.py`, and the figures of the best solutions are generated by `scripts/plot.py`.